

Common Search Commands

Command	Description
<code>chart/ timechart</code>	Returns results in a tabular output for (time-series) charting.
<code>dedup</code>	Removes subsequent results that match a specified criterion.
<code>eval</code>	Calculates an expression. See COMMON EVAL FUNCTIONS.
<code>fields</code>	Removes fields from search results.
<code>head/tail</code>	Returns the first/last N results.
<code>lookup</code>	Adds field values from an external source.
<code>rename</code>	Renames a field. Use wildcards to specify multiple fields.
<code>rex</code>	Specifies regular expression named groups to extract fields.
<code>search</code>	Filters results to those that match the search expression.
<code>sort</code>	Sorts the search results by the specified fields.
<code>stats</code>	Provides statistics, grouped optionally by fields. See COMMON STATS FUNCTIONS.
<code>mstats</code>	Similar to stats but used on metrics instead of events.
<code>table</code>	Specifies fields to keep in the result set. Retains data in tabular format.
<code>top/rare</code>	Displays the most/least common values of a field.
<code>transaction</code>	Groups search results into transactions.
<code>where</code>	Filters search results using eval expressions. Used to compare two different fields.

Common Eval Functions

The eval command calculates an expression and puts the resulting value into a field (e.g. "...| eval force = mass * acceleration"). The following table lists some of the functions used with the eval command. You can also use basic arithmetic operators (+ - * / %), string concatenation (e.g., "...| eval name = last . \" . first\"), and Boolean operations (AND OR NOT XOR < > <= >= != == LIKE).

Function	Description	Examples
abs (X)	Returns the absolute value of X.	abs (number)
case (X, "Y" ,...)	Takes pairs of arguments X and Y, where X arguments are Boolean expressions. When evaluated to TRUE, the arguments return the corresponding Y argument.	case(error == 404, "Not found", error == 500,"Internal Server Error", error == 200, "OK")
ceil (X)	Ceiling of a number X.	ceil(1.9)
cidrmatch ("X" , Y)	Identifies IP addresses that belong to a particular subnet.	cidrmatch("123.132.32.0/25",ip)
coalesce (X,...)	Returns the first value that is not null.	coalesce(null(), "Returned val", null())
cos (X)	Calculates the cosine of X.	n=cos(0)
exact (X)	Evaluates an expression X using double precision floating point arithmetic.	exact(3.14*num)
exp (X)	Returns eX.	exp(3)
if (X,Y,Z)	If X evaluates to TRUE, the result is the second argument Y. If X evaluates to FALSE, the result evaluates to the third argument Z.	if(error==200, "OK", "Error")
in (field,value-list)	Returns TRUE if a value in "value-list" matches a value in "field". You must use the "in" function inside the "if" function.	if(in(status, "404","500","503"),"true","false")
isbool (X)	Returns TRUE if X is Boolean.	isbool(field)
isint (X)	Returns TRUE if X is an integer.	isint(field)
isnull (X)	Returns TRUE if X is NULL.	isnull(field)
isstr ()	Returns TRUE if X is a string.	isstr(field)
len (X)	This function returns the character length of a string X.	len(field)
like (X, "Y")	Returns TRUE if and only if X is like the SQLite pattern in Y.	like(field, "addr%")
log (X,Y)	Returns the log of the first argument X using the second argument Y as the base. Y defaults to 10.	log(number,2)
lower (X)	Returns the lowercase of X.	lower(username)
ltrim (X,Y)	Returns X with the characters in Y trimmed from the left side. Y defaults to spaces and tabs.	ltrim(" ZZZabcZZ ", " Z")
match (X,Y)	Returns if X matches the regex pattern Y.	match(field, "^\\d{1,3}\\\\.\\d\$")
max (X,...)	Returns the maximum.	max(delay, mydelay)
md5 (X)	Returns the MD5 hash of a string value X.	md5(field)
min (X,...)	Returns the minimum.	min(delay, mydelay)
mvcount (X)	Returns the number of values of X.	mvcount(multifield)
mvfilter (X)	Filters a multi-valued field based on the Boolean expression X.	mvfilter(match(email, "net\$"))

mvindex (X,Y,Z)	Returns a subset of the multivalued field X from start position (zero-based) Y to Z (optional).	mvindex(multifield, 2)
mvjoin (X,Y)	Given a multi-valued field X and string delimiter Y, and joins the individual values of X using Y.	mvjoin(address, ";")
now ()	Returns the current time, represented in Unix time.	now ()
null ()	This function takes no arguments and returns NULL.	null ()
nullif (X,Y)	Given two arguments, fields X and Y, and returns the X if the arguments are different. Otherwise returns NULL.	nullif(fieldA, fieldB)
random ()	Returns a pseudo-random number ranging from 0 to 2147483647.	random ()
relative_time (X,Y)	Given epochtime time X and relative time specifier Y, returns the epochtime value of Y applied to X.	relative_time(now(), "-1d@d")
replace (X,Y,Z)	Returns a string formed by substituting string Z for every occurrence of regex string Y in string X.	Returns date with the month and day numbers switched, so if the input was 4/30/2021 the return value would be 30/4/2021: replace(date, " [^] (\\d{1,2})/(\\d{1,2})/", "\\2/\\1/")

Common Eval Functions (continued)

Function	Description	Examples
round (X, Y)	Returns X rounded to the amount of decimal places specified by Y. The default is to round to an integer.	<code>round(3.5)</code>
rtrim (X, Y)	Returns X with the characters in Y trimmed from the right side. If Y is not specified, spaces and tabs are trimmed.	<code>rtrim(" ZZZZabcZZ ", " Z")</code>
split (X, "Y")	Returns X as a multi-valued field, split by delimiter Y.	<code>split(address, ";")</code>
sqrt (X)	Returns the square root of X.	<code>sqrt(9)</code>
strftime (X, Y)	Returns epochtime value X rendered using the format specified by Y.	<code>strftime(_time, "%H:%M")</code>
strptime (X, Y)	Given a time represented by a string X, returns value parsed from format Y.	<code>strptime(timeStr, "%H:%M")</code>
substr (X, Y, Z)	Returns a substring field X from start position (1-based) Y for Z (optional) characters.	<code>substr("string", 1, 3)</code>
time ()	Returns the wall-clock time with microsecond resolution.	<code>time()</code>
tonumber (X, Y)	Converts input string X to a number, where Y (optional, defaults to 10) defines the base of the number to convert to.	<code>tonumber("0A4", 16)</code>
tostring (X, Y)	Returns a field value of X as a string. If the value of X is a number, it reformats it as a string. If X is a Boolean value,, reformats to "True" or "False". If X is a number, the second argument Y is optional and can either be "hex" (convert X to hexadecimal), "commas" (formats X with commas and 2 decimal places), or "duration" (converts seconds X to readable time format HH:MM:SS).	This example returns: <code>foo=615</code> and <code>foo2=00:10:15</code> : ... eval foo=615 eval foo2 = tostring(foo, "duration")
typeof (X)	Returns a string representation of the field type.	This example returns: <code>"NumberStringBoolInvalid": typeof(12)+ typeof("string")+</code>
urldecode (X)	Returns the URL X decoded.	<code>urldecode("http%3A%2F%2Fwww.splunk.com%2Fdownload%3Fr%3Dheader")</code>
validate (X, Y, ...)	Given pairs of arguments, Boolean expressions X and strings Y, returns the string Y corresponding to the first expression X that evaluates to False and defaults to NULL if all are True.	<code>validate(isint(port), "ERROR: Port is not an integer", port >= 1 AND port <= 65535, "ERROR: Port is out of range")</code>

Common Stats Functions

Common statistical functions used with the chart, stats, and timechart commands. Field names can be wildcarded, so `avg(*delay)` might calculate the average of the delay and `xdelay` fields.

avg (X)	Returns the average of the values of field X.
count (X)	Returns the number of occurrences of the field X. To indicate a specific field value to match, format X as <code>eval(field="value")</code> .
dc (X)	Returns the count of distinct values of the field X.
earliest (X)	Returns the chronologically earliest seen value of X.
latest (X)	Returns the chronologically latest seen value of X.
max (X)	Returns the maximum value of the field X. If the values of X are non-numeric, the max is found from alphabetical ordering.
median (X)	Returns the middle-most value of the field X.
min (X)	Returns the minimum value of the field X. If the values of X are non-numeric, the min is found from alphabetical ordering.
mode (X)	Returns the most frequent value of the field X.
perc<X> (Y)	Returns the X-th percentile value of the field Y. For example, <code>perc5(total)</code> returns the 5th percentile value of a field "total".
range (X)	Returns the difference between the max and min values of the field X.

stdev (X)	Returns the sample standard deviation of the field X.
stdevp (X)	Returns the population standard deviation of the field X.
sum (X)	Returns the sum of the values of the field X.
sumsq (X)	Returns the sum of the squares of the values of the field X.
values (X)	Returns the list of all distinct values of the field X as a multi-value entry. The order of the values is alphabetical.
var (X)	Returns the sample variance of the field X.

Search Examples

Filter Results	
Returns X rounded to the amount of decimal places specified by Y. The default is to round to an integer.	<code>round(3.5)</code>
Returns X with the characters in Y trimmed from the right side. If Y is not specified, spaces and tabs are trimmed.	<code>rtrim(" ZZZZabcZZ ", "Z")</code>
Returns X as a multi-valued field, split by delimiter Y.	<code>split(address, ";")</code>
Given pairs of arguments, Boolean expressions X and strings Y, returns the string Y corresponding to the first expression X that evaluates to False and defaults to NULL if all are True.	<code>validate(isint(port), "ERROR: Port is not an integer", port >= 1 AND port <= 65535, "ERROR: Port is out of range")</code>

Group Results	
Cluster results together, sort by their "cluster_count" values, and then return the 20 largest clusters (in data size).	<code>... cluster t=0.9 showcount=true sort limit=20 -cluster_count</code>
Group results that have the same "host" and "cookie", occur within 30 seconds of each other, and do not have a pause greater than 5 seconds between each event into a transaction.	<code>... transaction host cookie maxspan=30s maxpause=5s</code>
Group results with the same IP address (clientip) and where the first result contains "signon", and the last result contains "purchase".	<code>... transaction clientip startswith="signon" endswith="purchase"</code>

Order Results	
Return the first 20 results.	<code>... head 20</code>
Reverse the order of a result set.	<code>... reverse</code>
Sort results by "ip" value (in ascending order) and then by "url" value (in descending order).	<code>... sort ip, -url</code>
Return the last 20 results in reverse order.	<code>... tail 20</code>

Reporting (cont.)	
Return the average for each hour, of any unique field that ends with the string "lay" (e.g., delay, xdelay, relay, etc).	<code>... stats avg(*lay) by date_hour</code>
Return the 20 most common values of the "url" field.	<code>... top limit=20 url</code>
Return the least common values of the "url" field.	<code>... rare url</code>

Advanced Reporting	
Compute the overall average duration and add 'avgdur' as a new field to each event where the 'duration' field exists	<code>... eventstats avg(duration) as avgdur</code>
Find the cumulative sum of bytes.	<code>... streamstats sum(bytes) as bytes_total timechart max(bytes_total)</code>
Find anomalies in the field 'Close_Price' during the last 10 years.	<code>sourcetype=nasdaq earliest=-10y anomalydetection Close_Price</code>
Create a chart showing the count of events with a predicted value and range added to each event in the time-series.	<code>... timechart count predict count</code>
Computes a five event simple moving average for field 'count' and write to new field 'smoothed_count.'	<code>"... timechart count trendline sma5(count) as smoothed_count"</code>

Metrics	
List all of the metric names in the "_metrics" metric index.	<code> mcatalog values(metric_name) WHERE index=_metrics</code>
See examples of the metric data points stored in the "_metrics" metric index.	<code> mpreview index=_metrics target_per_timeseries=5</code>
Return the average value of a metric in the "_metrics" metric index. Bucket the results into 30 second time spans.	<code> mstats avg(aws.ec2.CPUUtilization) WHERE index=_metrics span=30s</code>

Reporting	
Return the average and count using a 30 second span of all metrics ending in cpu.percent split by each metric name.	<pre> mstats avg(_value), count(_value) WHERE metric_name="*.cpu. percent" by metric_name span=30s</pre>
Return max(delay) for each value of foo split by the value of bar.	<pre>... chart max(delay) over foo by bar</pre>
Return max(delay) for each value of foo.	<pre>... chart max(delay) over fo o</pre>
Count the events by "host"	<pre>... stats count by host</pre>
Create a table showing the count of events and a small line chart	<pre>... stats sparkline count by host</pre>
Create a timechart of the count of from "web" sources by "host"	<pre>... timechart count by host</pre>
Calculate the average value of "CPU" each minute for each "host".	<pre>... timechart span=1m avg(CPU) by host</pre>

Add Fields	
Set velocity to distance / time.	<pre>... eval velocity=distance/time</pre>
Extract "from" and "to" fields using regular expressions. If a raw event contains "From: Susan To: David", then from=Susan and to=David.	<pre>... rex field=_raw "From: (?<from>.*) To: (?<to>.*) "</pre>
Save the running total of "count" in a field called "total_count".	<pre>... accum count as total_count</pre>
For each event where 'count' exists, compute the difference between count and its previous value and store the result in 'countdiff'.	<pre>... delta count as countdiff</pre>

Filter Fields	
Keep only the "host" and "ip" fields, and display them in that order.	<pre>... fields + host, ip</pre>
Remove the "host" and "ip" fields from the results.	<pre>... fields - host, ip</pre>

Search Examples (continued)

Lookup Tables (Splunk Enterprise only)	
For each event, use the lookup table usertogroup to locate the matching "user" value from the event. Output the group field value to the event	<pre>... lookup usertogroup user output group</pre>
Read in the usertogroup lookup table that is defined in the transforms.conf file.	<pre>... inputlookup usertogroup</pre>
Write the search results to the lookup file "users.csv".	<pre>... outputlookup users.csv</pre>

Modify Fields	
Rename the "_ip" field as "IPAddress".	<pre>... rename _ip as IPAddress</pre>

Regular Expressions (Regexes)	
Regular Expressions are useful in multiple areas: search commands regex and rex; eval functions match() and replace(); and in field extraction.	

Multi-Valued Fields	
Combine the multiple values of the recipients field into a single value	<pre>... nomv recipients</pre>
Separate the values of the "recipients" field into multiple field values, displaying the top recipients	<pre>... makemv delim="," recipients top recipients</pre>
Create new results for each value of the multivalue field "recipients"	<pre>... mvexpand recipients</pre>
Find the number of recipient values	<pre>... eval to_count = mvcount(recipients)</pre>
Find the first email address in the recipient field	<pre>... eval recipient_first = mvindex(recipient,0)</pre>
Find all recipient values that end in .net or .org	<pre>... eval netorg_recipients = mvfilter match(recipient,"\.net\$") OR match(recipient,"\.org\$"))</pre>
Find the index of the first recipient value match "\.org\$"	<pre>... eval orgindex = mvfind(recipient, "\.org\$")</pre>

Regex	Note	Example	Explanation
\s	white space	\d\s\d	digit space digit
\S	not white space	\d\S\d	digit nonwhitespace digit
\d	digit	\d\d\d-\d\d-\d\d\d\d	SSN
\D	not digit	\D\D\D	three non-digits
\w	word character (letter, number, or _)	\w\w\w	three word chars
\W	not a word character	\W\W\W	three non-word chars
[...]	any included character	[a-z0-9#]	any char that is a thru z, 0 thru 9, or #
[^...]	no included character	[^xyz]	any char but x, y, or z
*	zero or more	\w*	zero or more words chars
+	one or more	\d+	integer
?	zero or one	\d\d\d-?\d\d-?\d\d\d\d	SSN with dashes being optional
	or	\w \d	word or digit character
(?P<var>...)	named extraction	(?P<ssn>\d\d\d-\d\d-\d\d\d\d)	pull out a SSN and assign to 'ssn' field
(?:...)	logical or atomic grouping	(?:[a-zA-Z] \d)	alphabetic character OR a digit
^	start of line	^\d+	line begins with at least one digit
\$	end of line	\d+\$	line ends with at least one digit
{...}	number of repetitions	\d{3,5}	between 3-5 digits
\	escape	\[escape the [character

Common Date and Time Formatting		
Use these values for eval functions strftime() and strptime(), and for timestamping event data.		
Time	%H	24 hour (leading zeros) (00 to 23)
	%I	12 hour (leading zeros) (01 to 12)
	%M	Minute (00 to 59)
	%S	Second (00 to 61)
	%N	subseconds with width (%3N = millisecs, %6N = microsecs, %9N = nanosecs)
	%p	AM or PM
	%Z	Time zone (EST)
	%z	Time zone offset from UTC, in hour and minute: +hhmm or -hhmm. (-0500 for EST)
	%s	Seconds since 1/1/1970 (1308677092)
Days	%d	Day of month (leading zeros) (01 to 31)
	%j	Day of year (001 to 366)
	%w	Weekday (0 to 6)
	%a	Abbreviated weekday (Sun)
	%A	Weekday (Sunday)
Months	%b	Abbreviated month name (Jan)
	%B	Month name (January)
	%m	Month number (01 to 12)
Years	%y	Year without century (00 to 99)
	%Y	Year (2021)
Examples	%Y-%m-%d	2021-12-31
	%y-%m-%d	21-12-31
	%b %d, %Y	Jan 24, 2021
	%B %d, %Y	January 24, 2021
	q %d %b '%y = %Y-%m-%d	q 25 Feb '21 = 2021-02-25